

▼ Problem Statement 2

Implement the Image classification CNN model for classifying hand-written MNIST dataset by dividing the model into following 4 stages:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, Flatten, Dense, MaxPooling2D
from tensorflow.keras.optimizers import Adam
```

a. Loading and preprocessing the image data

```
# Load the training data from CSV file
train_data = pd.read_csv('mnist_train.csv')
x_train = train_data.drop('label', axis=1).values
y_train = train_data['label'].values

# Load the testing data from CSV file
test_data = pd.read_csv('mnist_test.csv')
x_test = test_data.drop('label', axis=1).values
y_test = test_data['label'].values

num_classes = 10

# Normalize pixel values to the range [0, 1]
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

# Flatten the images to a 1D array (for MNIST)
x_train = x_train.reshape((-1, 28, 28, 1))
x_test = x_test.reshape((-1, 28, 28, 1))

# Convert labels to one-hot encoding
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

b. Defining the model's architecture

```
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    MaxPooling2D(2),
    Flatten(),
    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])

# Display the model summary
model.summary()
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
=====		
conv2d_12 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_8 (MaxPooling2D)	(None, 13, 13, 32)	0
flatten_4 (Flatten)	(None, 5408)	0
dense_8 (Dense)	(None, 128)	692352
dense_9 (Dense)	(None, 10)	1290
=====		
Total params: 693962 (2.65 MB)		
Trainable params: 693962 (2.65 MB)		
Non-trainable params: 0 (0.00 Byte)		

c. Training the model

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
model.fit(x_train, y_train, epochs=5, batch_size=64, validation_data=(x_test, y_test))
```

```
Epoch 1/5
938/938 [=====] - 39s 41ms/step - loss: 0.1759 - accuracy: 0.9490 - val_loss: 0.0635 - val_accuracy: 0.978
Epoch 2/5
938/938 [=====] - 38s 40ms/step - loss: 0.0581 - accuracy: 0.9823 - val_loss: 0.0508 - val_accuracy: 0.982
Epoch 3/5
938/938 [=====] - 49s 53ms/step - loss: 0.0382 - accuracy: 0.9879 - val_loss: 0.0510 - val_accuracy: 0.982
Epoch 4/5
938/938 [=====] - 39s 42ms/step - loss: 0.0258 - accuracy: 0.9917 - val_loss: 0.0369 - val_accuracy: 0.987
Epoch 5/5
938/938 [=====] - 39s 42ms/step - loss: 0.0183 - accuracy: 0.9945 - val_loss: 0.0454 - val_accuracy: 0.986
<keras.src.callbacks.History at 0x7ce1c8065270>
```

d. Estimating the model's performance

```
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss: ', score[0])
print('Test accuracy: ', score[1])

Test loss: 0.045391954481601715
Test accuracy: 0.9860000014305115

import random

n = random.randint(0,9999)
plt.imshow(x_test[n].reshape(28, 28), cmap='gray')
predicted_value = model.predict(x_test)
print("Actual Number: ",np.argmax(y_test[n]))
print("Predicted Number: ", np.argmax(predicted_value[n]))

313/313 [=====] - 5s 14ms/step
Actual Number: 0
Predicted Number: 0
```

